

Hive Case Study

E-Commerce Sales Review

By,
Antara Chatterji
Satvik Yadav

PROBLEM STATEMENT :

Tech companies are trying to analyse the customer behaviour and gain insights about the product trends. This helps the companies to make the products easily available for the customers and thus increasing their sales. As a big data analyst, one needs to extract the data and gather insights from real time data of an e-commerce company.

Objective: We need to gather the insights from the clickstream data so we can extract insights from the customer behaviour.

IMPLEMENTATION PHASE:

The implementation phase can be divided into the following parts:

Copying the data set into the HDFS:

- Launch an EMR cluster that utilizes the Hive services, and
- Move the data from the S3 bucket into the HDFS

Creating the database and launching Hive queries on your EMR cluster:

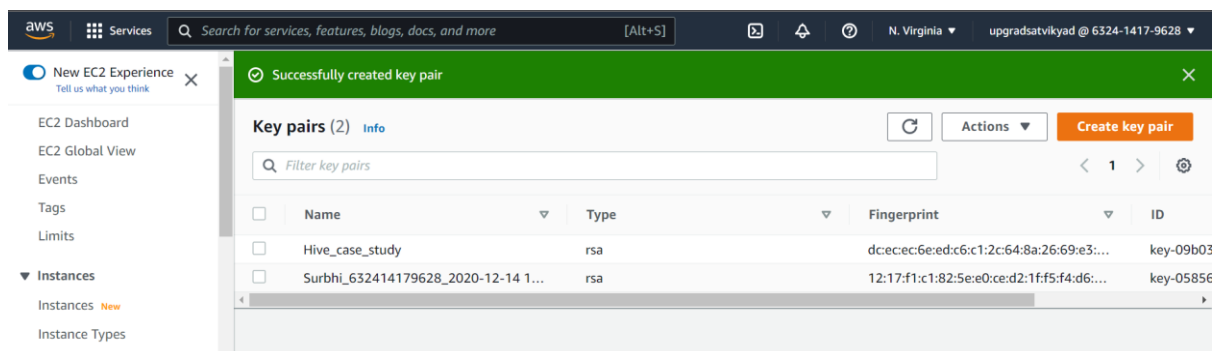
- Create the structure of your database,
- Use optimized techniques to run your queries as efficiently as possible
- Show the improvement of the performance after using optimization on any single query.
- Run Hive queries to answer the questions given below.

Cleaning up

- Drop your database, and
- Terminate your cluster

Launch the EMR cluster:

- To Launch the EMR cluster
- Create a key-pair and download the .PEM/.PPK file



- Now we need to create EMR cluster. While creating EMR cluster we need to select **m4.large** for both Master and Core node of single instance.

Cluster Nodes and Instances

Choose the instance type, number of instances, and a purchasing option. [Learn more about instance purchasing options](#)

Console options for automatic scaling have changed. [Learn more](#)

Node type	Instance type	Instance count	Purchasing option
Master Master - 1	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	On-demand Spot Use on-demand as max price
Core Core - 2	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	On-demand Spot Use on-demand as max price

+ Add task instance group

- Then we need to select the correct key pair from the dropdown.

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

Security Options

EC2 key pair: Hive_case_study

☒ Cluster visible to all IAM users in account

Permissions: ☒ Default ☐ Custom
Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role: EMR_DefaultRole ☐ Use EMR_DefaultRole_V2

EC2 instance profile: EMR_EC2_DefaultRole

Auto Scaling role: EMR_AutoScaling_DefaultRole

Security Configuration

EC2 security groups

- We need to provide a relevant cluster name while creating the cluster.

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

General Options

Cluster name: Hive_case_study_SY_AC

☒ Logging

S3 folder: s3://aws-logs-632414179628-us-east-1/elasticmap

☐ Log encryption


☒ Debugging

☒ Termination protection

Tags

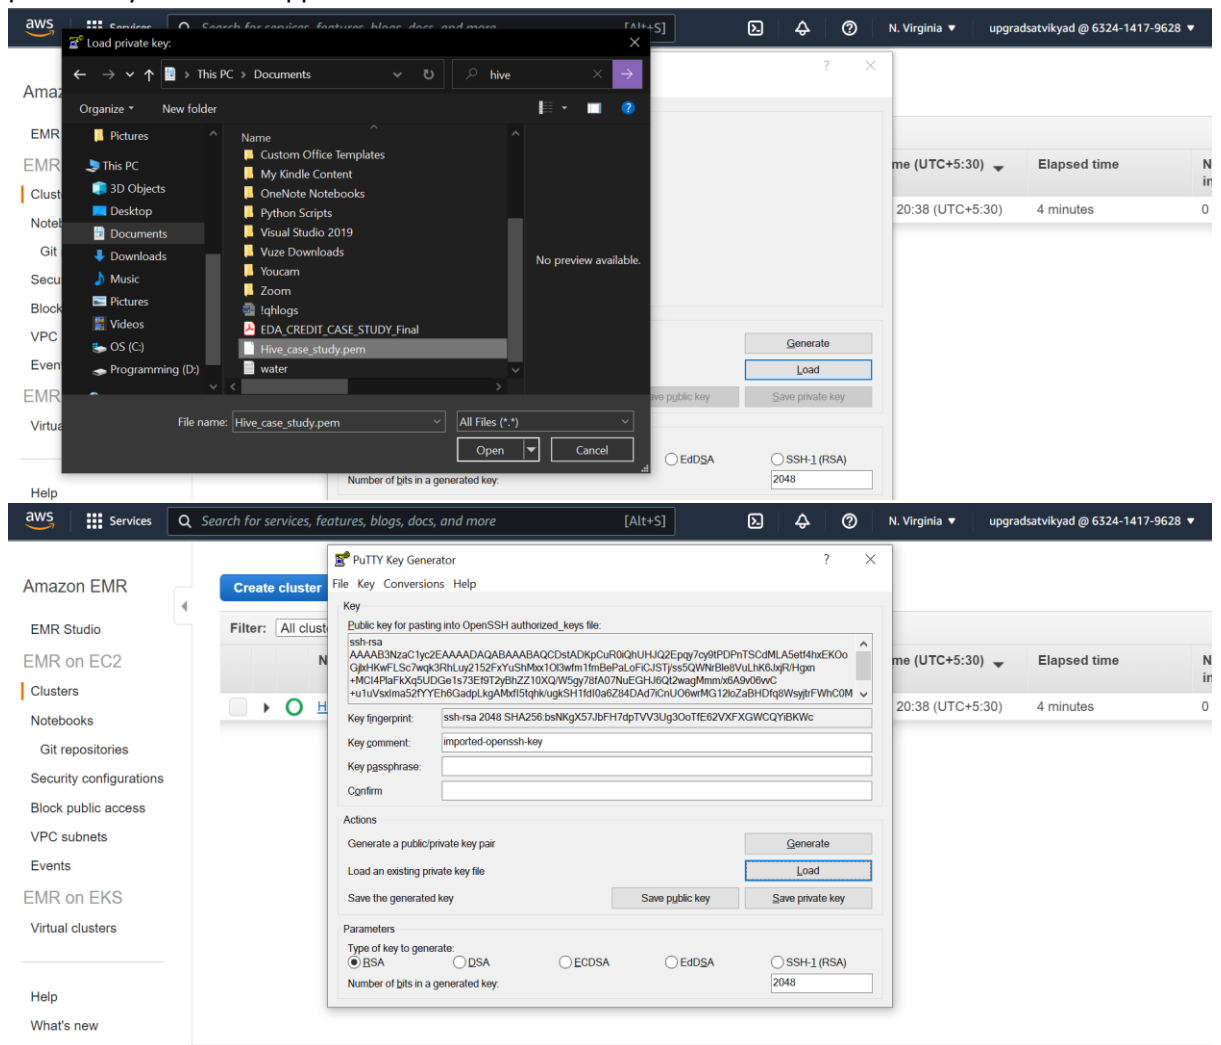
Key	Value (optional)
-----	------------------

- Now the cluster is created and is in the waiting stage. Now we need to move data from S3 to HDFS.

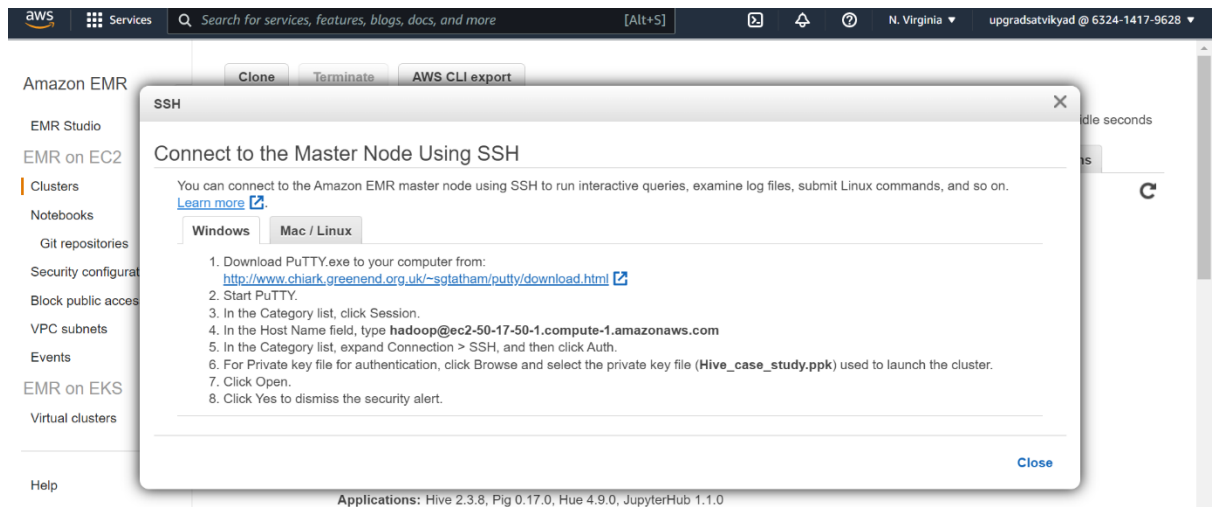


Connect to EMR cluster:

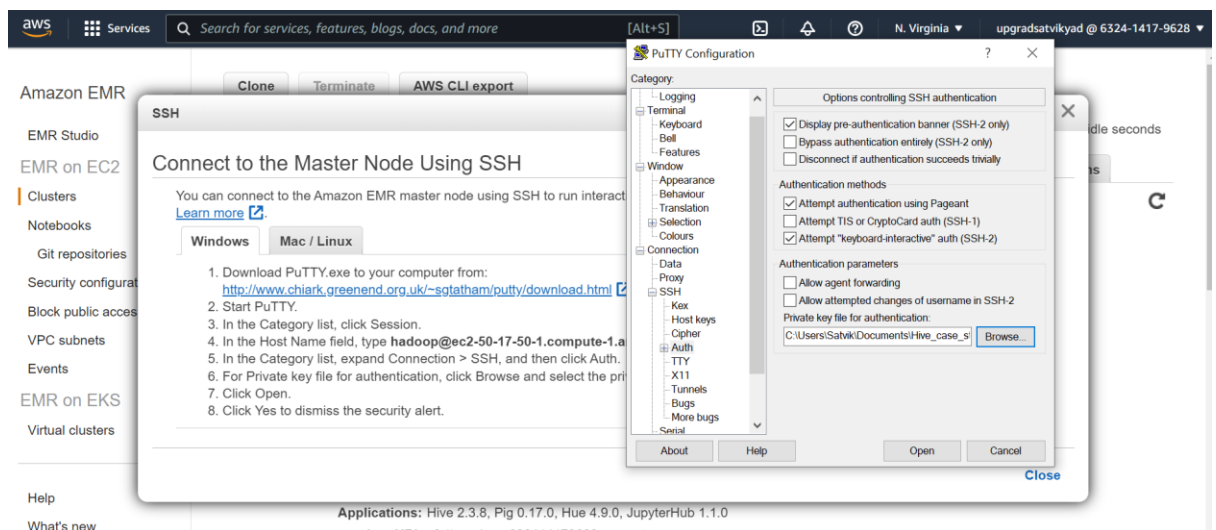
- We need to open PuTTY for windows and open the downloaded .pem file and save the private key which is in .ppk extension.



- When the cluster is in running state we click on Master public DNS.



- We need to open the PuTTY application and provide the Host Name(Master Node DNS) and then click on Connection>SSH>Auth, browse to the private key file location. Select the .ppk file and connect to the master node.



```
[hadoop@ip-172-31-20-27 ~]$ hadoop distcp s3://assignment-hive-datasets/2019-Oct
.csv /Hive_assignment/2019-Oct.csv
22/01/04 10:13:15 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=f
alse, syncFolder=false, deleteMissing=false, ignoreFailures=false, overwrite=fal
se, skipCRC=false, blocking=true, numListStatusThreads=0, maxMaps=20, mapBandwid
```

```

File Input Format Counters
  Bytes Read=227
File Output Format Counters
  Bytes Written=0
DistCp Counters
  Bytes Copied=482542278
  Bytes Expected=482542278
  Files Copied=1
[hadoop@ip-172-31-20-27 ~]$

```

View the data in HDFS by executing below commands:

Oct 2019 Sales data:

```

[hadoop@ip-172-31-20-27 ~]$ hadoop fs -cat /Hive_assignment/2019-Oct.csv | head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,u
ser_session
2019-10-01 00:00:00 UTC,cart,5773203,1487580005134238553,,runail,2.62,463240011,
26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:03 UTC,cart,5773353,1487580005134238553,,runail,2.62,463240011,
26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:07 UTC,cart,5881589,2151191071051219817,,lovely,13.48,429681830
,49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:07 UTC,cart,5723490,1487580005134238553,,runail,2.62,463240011,
26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:15 UTC,cart,5881449,1487580013522845895,,lovely,0.56,429681830,
49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:16 UTC,cart,5857269,1487580005134238553,,runail,2.62,430174032,
73deale7-664e-43f4-8b30-d32b9d5af04f
2019-10-01 00:00:19 UTC,cart,5739055,1487580008246412266,,kapous,4.75,377667011,
81326ac6-daa4-4f0a-b488-fd0956a78733
2019-10-01 00:00:24 UTC,cart,5825598,1487580009445982239,,0.56,467916806,2f5b55
46-b8cb-9ee7-7ecd-84276f8ef486
2019-10-01 00:00:25 UTC,cart,5698989,1487580006317032337,,1.27,385985999,d30965
e8-1101-44ab-b45d-cc1bb9fae694

```

Nov 2019 Sales Data:

```

[hadoop@ip-172-31-20-27 ~]$ hadoop fs -cat /Hive_assignment/2019-Nov.csv | head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,u
ser_session
2019-11-01 00:00:02 UTC,view,5802432,1487580009286598681,,,0.32,562076640,09fafd
6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC,cart,5844397,1487580006317032337,,,2.38,553329724,206721
6c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC,view,5837166,1783999064103190764,,pnb,22.22,556138645,57
ed222e-a54a-4907-9944-5a875c2d7f4f
2019-11-01 00:00:11 UTC,cart,5876812,1487580010100293687,,jessnail,3.16,56450666
6,186c1951-8052-4b37-adce-dd9644b1d5f7
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,,3.33,5533
29724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,,3.33,5533
29724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:25 UTC,view,5856189,1487580009026551821,,runail,15.71,562076640
,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:32 UTC,view,5837835,1933472286753424063,,,3.49,514649199,432a4e
95-375c-4b40-bd36-0fc039e77580
2019-11-01 00:00:34 UTC,remove_from_cart,5870838,1487580007675986893,,milv,0.79,
429913900,2f0bfff3c-252f-4fe6-afcd-5d8a6a92839a

```

After successfully adding the data, it's time to Set the data in Hive - Launch Hive :

```

[hadoop@ip-172-31-20-27 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.
properties Async: false
hive> create database if not exists Sales_Case_Study;
OK
Time taken: 1.761 seconds

```

Creating database, creating tables :

```
hive> describe database Sales_Case_Study;
OK
sales_case_study          hdfs://ip-172-31-20-27.ec2.internal:8020/user/hive/warehouse/sales_case_study.db      hadoop  USER
Time taken: 0.277 seconds, Fetched: 1 row(s)
hive> use Sales_Case_Study;
OK
Time taken: 0.037 seconds
hive>
```

Creating a table from the raw data given by taking care of the data dictionary:

```
[hadoop@ip-172-31-20-27 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Sales (
  > event_time timestamp,
  > event_type string,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string )
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
  > STORED AS TEXTFILE
  > LOCATION '/Hive_assignment'
  > TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 1.916 seconds
hive>
```

The Table schema:

```
hive> DESCRIBE Sales;
OK
event_time          string          from deserializer
event_type           string          from deserializer
product_id           string          from deserializer
category_id          string          from deserializer
category_code        string          from deserializer
brand                string          from deserializer
price                string          from deserializer
user_id              string          from deserializer
user_session         string          from deserializer
Time taken: 0.493 seconds, Fetched: 9 row(s)
hive>
```

Loading data from both the files into this table:

```
hive> LOAD DATA INPATH '/Hive_assignment/2019-Oct.csv' into table Sales;
Loading data to table default.sales
OK
Time taken: 1.03 seconds
hive> LOAD DATA INPATH '/Hive_assignment/2019-Nov.csv' into table Sales;
Loading data to table default.sales
OK
Time taken: 0.533 seconds
hive>
```


Viewing the data after loading them into the table:

```
hive> SELECT * FROM Sales LIMIT 3;
OK
2019-11-01 00:00:02 UTC view      5802432 1487580009286598681      0.
32      562076640      09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC cart      5844397 1487580006317032337      2.
38      553329724      2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC view      5837166 1783999064103190764      pnb      22
.22      556138645      57ed222e-a54a-4907-9944-5a875c2d7f4f
Time taken: 1.911 seconds, Fetched: 3 row(s)
hive>
```

Creating another table for data analysis with data in proper format.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Sales_Data (
> event_time timestamp,
> event_type string,
> product_id string,
> category_id string,
> brand string,
> price float,
> user_id bigint,
> user_session string )
> ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
> LINES TERMINATED BY '\n'
> STORED AS TEXTFILE;
OK
Time taken: 0.039 seconds
hive>
```

Inserting the data into this table (Sales_Data).

```
hive> INSERT INTO Sales_data
> SELECT
> cast(replace(event_time,'UTC','') as timestamp),
> event_type,
> product_id,
> category_id,
> category_code,
> brand,
> cast(price as float),
> cast(user_id as bigint),
> user_session
> from Ecom;
Query ID = hadoop_20220104103620_a740a0aa-5420-4962-8f01-4879a91b08d6
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1641290350049
```

Data loaded successfully:

```
Map 1: 1(+1)/2
Map 1: 2/2
Loading data to table default.sales_data
OK
Time taken: 150.641 seconds
hive>
```


Running the first query in order to check the total time taken by the query to run successfully without partitioning:

Q1. Find the total revenue generated due to purchases made in October. (using the Sales_data table)

```
hive> SELECT SUM(price) As total_revenue
> FROM Sales_data
> WHERE MONTH(event_time)=10 AND event_type='purchase';
Query ID = hadoop_20220104124149_28beaa37-67f0-45ea-b91e-569414fe7d56
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

Map 1: 0/6      Reducer 2: 0/1
Map 1: 0/6      Reducer 2: 0/1
Map 1: 0/6      Reducer 2: 0/1
Map 1: 0(+1)/6  Reducer 2: 0/1
Map 1: 0(+2)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 0(+3)/6  Reducer 2: 0/1
Map 1: 1(+2)/6  Reducer 2: 0/1
Map 1: 1(+3)/6  Reducer 2: 0/1
Map 1: 2(+3)/6  Reducer 2: 0/1
Map 1: 3(+2)/6  Reducer 2: 0/1
Map 1: 3(+3)/6  Reducer 2: 0/1
Map 1: 3(+3)/6  Reducer 2: 0/1
Map 1: 3(+3)/6  Reducer 2: 0/1
Map 1: 4(+2)/6  Reducer 2: 0/1
Map 1: 4(+2)/6  Reducer 2: 0(+1)/1
Map 1: 5(+1)/6  Reducer 2: 0(+1)/1
Map 1: 6/6      Reducer 2: 0(+1)/1
Map 1: 6/6      Reducer 2: 1/1
OK
NULL
Time taken: 37.564 seconds, Fetched: 1 row(s)
hive>
```

SELECT SUM(price) AS total_revenue

FROM Sales_data

WHERE MONTH(event_time)=10 AND event_type='purchase';

**Time Taken by the entire query without table optimization (i.e. without partitioning) :
37.564 seconds**

Note: Once the base table is created, we need to optimize the table for quick query result through partitioning and bucketing. Our first optimized table name is Sales_data_part1.

Now we will be enabling Dynamic Partitioning and creating a partitioned table with buckets.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Sales_Data_Part1 (
  > event_time timestamp,
  > event_type string,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string )
  > PARTITIONED BY (year int, month int)
  > CLUSTERED BY (category_id) into 4 Buckets
  > ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
  > LINES TERMINATED BY '\n'
  > STORED AS TEXTFILE;
OK
Time taken: 0.071 seconds
hive>
```

```
OK
event_time          timestamp
event_type          string
product_id          string
category_id         string
category_code       string
brand               string
price              float
user_id             bigint
user_session       string
year                int
month               int

# Partition Information
# col_name          data_type          comment

year                int
month               int
Time taken: 0.688 seconds, Fetched: 17 row(s)
hive>
```

The new optimised table has been created. We will now insert the data into this table.

```
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> INSERT INTO table Sales_data_part1 Partition (year,month)
  > select
  > cast(replace(event_time,'UTC','') as timestamp),
  > event_type,
  > product_id,
  > category_id,
  > category_code,
  > brand,
  > cast(Price as float),
  > cast(user_id as bigint),
  > user_session,
  > year(cast(replace(event_time,'UTC','') as timestamp)),
  > month(cast(replace(event_time,'UTC','') as timestamp))
  > from Sales;
Query ID = hadoop_20220104124444_f2183579-d933-4052-a4f1-7cc613caf38d
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 02/02 [=====>] 100% ELAPSED TIME: 222.51 s
-----
Loading data to table default.sales_data_part1 partition (year=null, month=null)

Loaded : 2/2 partitions.
Time taken to load dynamic partitions: 0.264 seconds
Time taken for adding to write entity : 0.0 seconds
OK
Time taken: 224.016 seconds
hive>
```

We will now test this table by running the same query (Q1) in this optimised table and note the time taken.

Q1. Find the total revenue generated due to purchases made in October. (using the Sales_data_part1 table)

```
hive> SELECT SUM(price)
> FROM Sales_Data_Part1
> WHERE month(event_time)=10 AND event_type = 'purchase';
Query ID = hadoop_20220104124846_bc552a33-ba2a-4639-b647-fce186b538c2
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container    SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100%  ELAPSED TIME: 36.92 s
-----
OK
1211538,4295325726
Time taken: 37.649 seconds, Fetched: 1 row(s)
hive>
```

SELECT SUM(price) AS total_revenue

FROM Sales_data_part1

WHERE MONTH(event_time)=10 AND event_type='purchase';

Time Taken by the entire query after dynamic partitioning of the table : 37.649 seconds

Enabling second approach for Dynamic Partitioning and creating another partitioned table with buckets.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Sales_Data_Part2 (
> event_time timestamp,
> product_id string,
> category_id string,
> category_code string,
> brand string,
> price float,
> user_id bigint,
> user_session string )
> PARTITIONED BY (event_type string)
> CLUSTERED BY (category_id) into 5 buckets
> ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
> STORED as TEXTFILE;
OK
Time taken: 0.078 seconds
hive>
```

The Second Partitioned table has been created . This table is named as Sales_Data_Part2.

```
hive> DESCRIBE Sales_Data_Part2;
OK
event_time          string          from deserializer
product_id          string          from deserializer
category_id         string          from deserializer
category_code       string          from deserializer
brand               string          from deserializer
price               string          from deserializer
user_id             string          from deserializer
user_session        string          from deserializer
event_type          string

# Partition Information
# col_name          data_type      comment

event_type          string
Time taken: 0.083 seconds, Fetched: 14 row(s)
hive>
```

We will now insert the data into this table:

```
hive> INSERT INTO table Sales_data_part2 Partition (event_type)
> select
> cast(replace(event_time,'UTC','') as timestamp),
> product_id,
> category_id,
> category_code,
> brand,
> cast(price as float),
> cast(user_id as bigint),
> user_session,
> event_type
> from Sales;
Query ID = hadoop_20220104125046_4ac2102f-dcff-4a84-9159-521aaadbb032
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 02/02 [=====>>] 100%  ELAPSED TIME: 217.11 s
-----
Loading data to table default.sales_data_part2 partition (event_type=null)

Loaded : 4/4 partitions.
    Time taken to load dynamic partitions: 0.375 seconds
    Time taken for adding to write entity : 0.001 seconds
OK
Time taken: 218.811 seconds
hive>
```

We will now test this table by running the same query in this optimised table and note the time taken.

Q1. Find the total revenue generated due to purchases made in October. (using the Sales_Data_Part2 table)

```
SELECT SUM(price) AS total_revenue

FROM Sales_data_Part2

WHERE MONTH(event_time)=10 AND event_type='purchase';
```

Time Taken by the entire query after dynamic partitioning of the table : 27.992 seconds

```
hive>
>
> SELECT SUM(price) As total_revenue
> FROM Sales_data_part2
> WHERE MONTH(event_time)=10 AND event_type='purchase';
Query ID = hadoop_20220104125446_a26a4d47-f51f-4952-8f82-66066e773b63
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3          3          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 26.97 s
-----
OK
1211538.4299998898
Time taken: 27.992 seconds, Fetched: 1 row(s)
hive>
```

CONCLUSION : It can be observed that, by Partition by over 'event_type' and clustering by 'category_id' we get the most optimized output of the query.

Note: The following questions are being solved using the most optimized table i.e. Sales_Data_Part2:

Q1. Find the total revenue generated due to purchases made in October.

```
hive>
>
> SELECT SUM(price) As total_revenue
> FROM Sales_data_part2
> WHERE MONTH(event_time)=10 AND event_type='purchase';
Query ID = hadoop_20220104125446_a26a4d47-f51f-4952-8f82-66066e773b63
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3          3          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 26.97 s
-----
OK
1211538.4299998898
Time taken: 27.992 seconds, Fetched: 1 row(s)
hive>
```

```
SELECT SUM(price) As total_revenue
FROM Sales_data_part2
WHERE MONTH(event_time)=10 AND event_type='purchase';
```

OUTPUT : Total Revenue = 1211538.4299

Q2. Write a query to yield the total sum of purchases per month in a single output.

```
hive> select month (event_time), sum(price) from Sales_data_part2 where year (event_time)=2019
> and event_type='purchase' group by month(event_time);
Query ID = hadoop_20220104125545_8d93ec35-f24d-44dc-9728-0a2473106723
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100%  ELAPSED TIME: 28.13 s
-----
OK
10      1211538.4299998898
11      1531016.9
Time taken: 28.762 seconds, Fetched: 2 row(s)
hive>
```

SELECT month (event_time), sum(price)

FROM Sales_data_part2

WHERE year (event_time)=2019 **AND** event_type='purchase'

GROUP BY month(event_time);

OUTPUT:

event_time (month)	Total_Purchase
10	1211538.4299
11	1531016.9

Q3. Write a query to find the change in revenue generated due to purchases from October to November.

```
hive> select sum (case when month(event_time)=10 then price else -1*price end) as
> change_in_revenue from Sales_data_part2 where month(event_time) in (10,11) a
nd
> event_type='purchase';
Query ID = hadoop_20220105063830_24408271-0850-41fe-95ed-9a2aab736acf
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641362417569_0
005)

Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0/3      Reducer 2: 0/1
Map 1: 0(+2)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 0(+3)/3  Reducer 2: 0/1
Map 1: 1(+2)/3  Reducer 2: 0(+1)/1
Map 1: 2(+1)/3  Reducer 2: 0(+1)/1
Map 1: 3/3      Reducer 2: 1/1
OK
-319478.47000012523
Time taken: 27.802 seconds, Fetched: 1 row(s)
hive>
```

SELECT sum(case when month(event_time)=10 then price else -1*price end) as change_in_revenue
FROM Sales_data_part2

WHERE month(event_time) in (10,11) and event_type='purchase';

OUTPUT: -319478.4700

Q4. Find distinct categories of products. Categories with null category code can be ignored.

```
hive> select distinct split(category_code,'\\\.')[0] as cat from Sales_data_part2 where
> split(category_code,'\\\.')[0] <> '';
Query ID = hadoop_20220104125746_e8c88a22-ecf4-4ae0-8590-5474d199e1b9
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   6         6           0         0         0         0
Reducer 2 ..... container  SUCCEEDED   5         5           0         0         0         0
-----
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 73.26 s
-----
OK
furniture
appliances
accessories
apparel
sport
stationery
Time taken: 74.062 seconds, Fetched: 6 row(s)
hive>
```

SELECT DISTINCT split(category_code,'\\\.')[0] as cat

FROM Sales_data_part2

WHERE split(category_code,'\\\.')[0] <> '';

OUTPUT:

Categories of Product
Furniture
Appliances
Accessories
Apparel
Sport
Stationery

Q5. Find the total number of products available under each category.

```
hive> SELECT SPLIT(category_code,'\\.')[0] AS cat, COUNT(product_id) AS
> No_of_products FROM Sales_data_part2 WHERE SPLIT(category_code,'\\.')[0] <> ''
> GROUP BY SPLIT(category_code,'\\.')[0] ORDER BY No_of_products DESC;
Query ID = hadoop_20220105063911_5c0b4abd-e37b-475d-aa72-5c671606c4c1
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641362417569_0005)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED   6         6          0         0         0         0
Reducer 2 ..... container    SUCCEEDED   5         5          0         0         0         0
Reducer 3 ..... container    SUCCEEDED   1         1          0         0         0         0
-----
VERTICES: 03/03  [=====>>>] 100%  ELAPSED TIME: 70.37 s
-----
OK
appliances      61736
stationery      26722
furniture       23604
apparel 18232
accessories     12929
sport           2
Time taken: 71.131 seconds, Fetched: 6 row(s)
```

SELECT SPLIT(category_code,'\\.')[0] AS cat, COUNT(product_id) AS

No_of_products FROM Sales_data_part2 WHERE SPLIT(category_code,'\\.')[0] <> ''

GROUP BY SPLIT(category_code,'\\.')[0] ORDER BY No_of_products DESC;

OUTPUT:

appliances	61736
Stationary	26722
Furniture	23604
Apparel	18232
Accessories	12929
Sport	2

Q6. Which brand had the maximum sales in October and November combined?

```
hive> SELECT brand, SUM (price) AS sales FROM Sales_data_part2 WHERE BRAND <>' ' and
> event_type='purchase' GROUP BY brand ORDER BY sales DESC LIMIT 1;
Query ID = hadoop_20220104125917_bf95a2e3-7b56-46a8-b9f9-52c6bc6c9ea0
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED   3         3          0         0         0         0
Reducer 2 ..... container    SUCCEEDED   1         1          0         0         0         0
Reducer 3 ..... container    SUCCEEDED   1         1          0         0         0         0
-----
VERTICES: 03/03  [=====>>>] 100%  ELAPSED TIME: 24.69 s
-----
OK
runail 148297.9400000049
Time taken: 25.301 seconds, Fetched: 1 row(s)
hive>
```

SELECT brand, SUM (price) AS sales

FROM Sales_data_part2

WHERE BRAND <>' ' and event_type='purchase' GROUP BY brand ORDER BY sales DESC LIMIT 1;

OUTPUT:

Brand Name	Max_Sales
runall	148297.94

Q7. Which brands increased their sales from October to November?

```
hive> WITH Monthly_rev AS (
  > SELECT brand,
  > SUM(CASE WHEN date_format(event_time, 'MM')=10 THEN price ELSE 0 END) AS Oct_rev,
  > SUM(CASE WHEN date_format(event_time, 'MM')=11 THEN price ELSE 0 END) AS Nov_rev
  > FROM Sales_data_part2
  > WHERE event_type='purchase' AND date_format(event_time, 'MM') IN ('10', '11')
  > GROUP BY brand )
  > SELECT brand, Oct_rev, Nov_rev, Nov_rev-Oct_rev AS Sales_diff
  > FROM Monthly_rev
  > WHERE (Nov_rev - Oct_rev)>0
  > ORDER BY Sales_diff;
Query ID = hadoop_20220105064040_e54c58a3-8714-49c9-bc63-4b49457cdd92
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641362417569_0005)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3          0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1          0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1          0         0         0         0
-----
VERTICES: 03/03  [=====>>>] 100%  ELAPSED TIME: 29.39 s
-----
OK
ovale 2.54 3.1 0.56
cosima 20.23 20.930000000000003 0.7000000000000028
grace 100.920000000000002 102.61000000000004 1.6900000000000261
helloganic 0.0 3.1 3.1
skinity 8.88 12.440000000000001 3.5600000000000005
bodyton 1376.3400000000001 1380.64 4.299999999999545
moyou 5.71 10.280000000000001 4.570000000000001
neoleor 43.41 51.7 8.290000000000006
soleo 204.200000000000027 212.53000000000014 8.329999999999987

metzger 5373.449999999997 6457.1599999999935 1083.7099999999964
de.lux 1659.699999999784 2775.509999999972 1115.8099999999936
swarovski 1887.9299999999891 3043.159999999999 1155.23000000000007
beauty-free 554.17000000000003 1782.8599999999998 1228.6899999999978
zeitun 708.6599999999999 2009.6299999999997 1300.9699999999998
joico 705.52 2015.1 1309.58
severina 4775.879999999998 6120.4799999999988 1344.59999999999894
irisk 45591.9599999999046 46946.039999999928 1354.08000000002346
oniq 8425.4099999999976 9841.649999999978 1416.24000000000016
levrana 2243.56000000000004 3664.1000000000003 1420.54000000000027
roubloff 3491.3599999999965 4913.7700000000014 1422.41000000000176
smart 4457.2599999999992 5902.1400000000007 1444.88000000000147
shik 3341.20000000000007 4839.7200000000001 1498.52000000000004
domix 10472.0500000000021 12009.1700000000031 1537.120000000001
artex 2730.6399999999994 4327.2499999999997 1596.60999999999979
beautix 10493.949999999986 12222.949999999997 1729.0000000000011
milv 3904.9399999999983 5642.0099999999976 1737.0699999999993
masura 31266.079999999823 33058.4699999998706 1792.3900000000476
f.o.x 6624.23 8577.280000000001 1953.05000000000102
kapous 11927.1600000000113 14093.0800000000078 2165.9199999999655
concept 11032.139999999974 13380.399999999994 2348.2599999999657
estel 21756.7500000000084 24142.670000000007 2385.9199999999873
kaypro 881.34 3268.7000000000003 2387.36
benovy 409.6199999999999 3259.969999999992 2850.349999999992
italwax 21940.23999999973 24799.369999999766 2859.1300000000374
yoko 8756.909999999994 11707.879999999965 2950.9699999999702
haruyama 9390.690000000014 12352.909999999999 2962.219999999985
marathon 7280.749999999997 10273.099999999986 2992.3499999999885
lovely 8704.379999999999 11939.0600000000029 3234.6800000000385
bpw.style 11572.1500000001808 14837.4400000002425 3265.29000000006175
staleks 8519.7300000000023 11875.609999999999 3355.8799999999756
freedecor 3421.7799999999706 7671.8000000000216 4250.0200000000245
runail 71539.279999999619 76758.659999999736 5219.3800000001169
polarus 6013.72000000000075 11371.9300000000013 5358.21000000000055
cosmoprofi 8322.809999999996 14536.989999999958 6214.179999999962
jessnail 26287.840000000013 33345.230000000008 7057.389999999952
strong 29196.629999999994 38671.269999999994 9474.64
ingarden 23161.390000000004 33566.20999999995 10404.8199999999057
lianail 5892.839999999998 16394.2400000000194 10501.4000000000214
uno 35302.030000000014 51039.7499999998894 15737.7199999998757
grattol 35445.540000000078 71472.70999999995 36027.169999999872
474679.059999999656 619509.2399999989 144830.17999999993
Time taken: 29.992 seconds, Fetched: 161 row(s)
hive>
```

```

WITH Monthly_rev AS (

SELECT brand,

SUM(CASE WHEN date_format(event_time, 'MM')=10 THEN price ELSE 0 END) AS Oct_rev,

SUM(CASE WHEN date_format(event_time, 'MM')=11 THEN price ELSE 0 END) AS Nov_rev

FROM Sales_data_part2

WHERE event_type='purchase' AND date_format(event_time, 'MM') IN ('10', '11')

GROUP BY brand )

SELECT brand, Oct_rev, Nov_rev, Nov_rev-Oct_rev AS Sales_diff

FROM Monthly_rev

WHERE (Nov_rev - Oct_rev)>0

ORDER BY Sales_diff;

```

OUTPUT: Total 161 rows returned (refer to the output screen)

Q8. Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.

```

Time taken: 27.991 seconds, Fetched: 1 row(s)
hive> SELECT user_id, SUM(price) as Total_Expense
> FROM Sales_data_part2
> WHERE event_type='purchase'
> GROUP BY user_id
> ORDER BY Total_Expense DESC
> LIMIT 10;
Query ID = hadoop_20220104130009_50d774ab-c4fd-4df9-b2d1-9a08e06d2151
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1641290350049_0013)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   3         3           0         0         0         0
Reducer 2 ..... container  SUCCEEDED   1         1           0         0         0         0
Reducer 3 ..... container  SUCCEEDED   1         1           0         0         0         0
-----
VERTICES: 03/03  [=====>>] 100%  ELAPSED TIME: 27.18 s
-----
OK
557790271      2715.8699999999995
150318419      1645.97
562167663      1352.85
531900924      1329.4499999999998
557850743      1295.48
522130011      1185.3899999999999
561592095      1109.7000000000003
431950134      1097.5899999999997
566576008      1056.3600000000006
521347209      1040.91
Time taken: 27.995 seconds, Fetched: 10 row(s)
hive>

```

```

SELECT user_id, SUM(price) as Total_Expense

FROM Sales_data_part2

WHERE event_type='purchase'

```

GROUP BY user_id

ORDER BY Total_Expense DESC

LIMIT 10;

OUTPUT:

User_id	Total_Amt_Spent
557790271	2715.87
150318415	1645.97
562167663	1382.85
531900924	1329.4499
557850743	1295.48
522130011	1329.4499
561592095	1109.70000
431950134	1097.58999
566576008	1056.36000
521347209	1040.91